


Advanced Machine Learning


Support Vector Machines



Support Vector Machines, maximum margins

Outline

1. The Big Picture
2. Practice
3. Maximum Margins



Support Vector Machines

- Support vector machines, when used right, often have the best generalisation results
- They are typically used on numerical data, but can and have been adapted to text, sequences, etc.
- Although not as trendy as deep learning, they will often be the method of choice on small data sets
- They subtly regularise themselves, choosing a solution that generalises well from a host of different solutions

Support Vector Machines

- Support vector machines, when used right, often have the best generalisation results
- They are typically used on numerical data, but can and have been adapted to text, sequences, etc.
- Although not as trendy as deep learning, they will often be the method of choice on small data sets
- They subtly regularise themselves, choosing a solution that generalises well from a host of different solutions

Support Vector Machines


- Support vector machines, when used right, often have the best generalisation results
- They are typically used on numerical data, but can and have been adapted to text, sequences, etc.
- Although not as trendy as deep learning, they will often be the method of choice on small data sets
- They subtly regularise themselves, choosing a solution that generalises well from a host of different solutions

Support Vector Machines

- Support vector machines, when used right, often have the best generalisation results
- They are typically used on numerical data, but can and have been adapted to text, sequences, etc.
- Although not as trendy as deep learning, they will often be the method of choice on small data sets
- They subtly regularise themselves, choosing a solution that generalises well from a host of different solutions


Linear Separation of Data

- SVMs classify linearly separable data




Linear Separation of Data

- SVMs classify linearly separable data




Linear Separation of Data

- SVMs classify linearly separable data




Linear Separation of Data

- SVMs classify linearly separable data




Linear Separation of Data

- SVMs classify linearly separable data




Linear Separation of Data

- SVMs classify linearly separable data




Linear Separation of Data

- SVMs classify linearly separable data



Linear Separation of Data

- SVMs classify linearly separable data



- Finds maximum-margin separating plane

Extended Feature Space

- To increase the likelihood of linear-separability we often use a high-dimensional mapping

$$\mathbf{x} = (x_1, x_2, \dots, x_p) \rightarrow \vec{\phi}(\mathbf{x}) = (\phi_1(\mathbf{x}), \phi_2(\mathbf{x}), \dots, \phi_m(\mathbf{x}))$$

$$m \gg p$$

- Finding the maximum margin hyper-plane is time consuming in “primal” form if m is large
- We can work in the “dual” space of patterns, then we only need to compute dot products

$$\vec{\phi}(\mathbf{x}_i) \cdot \vec{\phi}(\mathbf{x}_j)$$

Extended Feature Space

- To increase the likelihood of linear-separability we often use a high-dimensional mapping

$$\mathbf{x} = (x_1, x_2, \dots, x_p) \rightarrow \vec{\phi}(\mathbf{x}) = (\phi_1(\mathbf{x}), \phi_2(\mathbf{x}), \dots, \phi_m(\mathbf{x}))$$

$$m \gg p$$

- Finding the maximum margin hyper-plane is time consuming in “primal” form if m is large
- We can work in the “dual” space of patterns, then we only need to compute dot products

$$\vec{\phi}(\mathbf{x}_i) \cdot \vec{\phi}(\mathbf{x}_j)$$

Extended Feature Space

- To increase the likelihood of linear-separability we often use a high-dimensional mapping

$$\mathbf{x} = (x_1, x_2, \dots, x_p) \rightarrow \vec{\phi}(\mathbf{x}) = (\phi_1(\mathbf{x}), \phi_2(\mathbf{x}), \dots, \phi_m(\mathbf{x}))$$

$$m \gg p$$

- Finding the maximum margin hyper-plane is time consuming in “primal” form if m is large
- We can work in the “dual” space of patterns, then we only need to compute dot products

$$\vec{\phi}(\mathbf{x}_i) \cdot \vec{\phi}(\mathbf{x}_j)$$

Extended Feature Space

- To increase the likelihood of linear-separability we often use a high-dimensional mapping

$$\mathbf{x} = (x_1, x_2, \dots, x_p) \rightarrow \vec{\phi}(\mathbf{x}) = (\phi_1(\mathbf{x}), \phi_2(\mathbf{x}), \dots, \phi_m(\mathbf{x}))$$

$$m \gg p$$

- Finding the maximum margin hyper-plane is time consuming in “primal” form if m is large
- We can work in the “dual” space of patterns, then we only need to compute dot products

$$\vec{\phi}(\mathbf{x}_i) \cdot \vec{\phi}(\mathbf{x}_j) = \sum_{k=1}^m \phi_k(\mathbf{x}_i) \phi_k(\mathbf{x}_j)$$

Kernel Trick

- If we choose a **positive semi-definite** kernel function $K(\mathbf{x}, \mathbf{y})$ then there exists functions $\phi_k(\mathbf{x})$, such that

$$K(\mathbf{x}_i, \mathbf{x}_j) = \vec{\phi}(\mathbf{x}_i) \cdot \vec{\phi}(\mathbf{x}_j)$$

(like an eigenvector decomposition of a matrix)

- Never need to compute $\phi_k(\mathbf{x}_i)$ explicitly as we only need the dot-product $\vec{\phi}(\mathbf{x}_i) \cdot \vec{\phi}(\mathbf{x}_j) = K(\mathbf{x}_i, \mathbf{x}_j)$ to compute maximum margin separating hyper-plane
- Sometimes $\vec{\phi}(\mathbf{x}_i)$ is an infinite dimensional vector so its good we don't have to compute it!

Kernel Trick

- If we choose a **positive semi-definite** kernel function $K(\mathbf{x}, \mathbf{y})$ then there exists functions $\phi_k(\mathbf{x})$, such that

$$K(\mathbf{x}_i, \mathbf{x}_j) = \vec{\phi}(\mathbf{x}_i) \cdot \vec{\phi}(\mathbf{x}_j)$$

(like an eigenvector decomposition of a matrix)

- Never need to compute $\phi_k(\mathbf{x}_i)$ explicitly as we only need the dot-product $\vec{\phi}(\mathbf{x}_i) \cdot \vec{\phi}(\mathbf{x}_j) = K(\mathbf{x}_i, \mathbf{x}_j)$ to compute maximum margin separating hyper-plane
- Sometimes $\vec{\phi}(\mathbf{x}_i)$ is an infinite dimensional vector so its good we don't have to compute it!

Kernel Trick

- If we choose a **positive semi-definite** kernel function $K(\mathbf{x}, \mathbf{y})$ then there exists functions $\phi_k(\mathbf{x})$, such that

$$K(\mathbf{x}_i, \mathbf{x}_j) = \vec{\phi}(\mathbf{x}_i) \cdot \vec{\phi}(\mathbf{x}_j)$$

(like an eigenvector decomposition of a matrix)

- Never need to compute $\phi_k(\mathbf{x}_i)$ explicitly as we only need the dot-product $\vec{\phi}(\mathbf{x}_i) \cdot \vec{\phi}(\mathbf{x}_j) = K(\mathbf{x}_i, \mathbf{x}_j)$ to compute maximum margin separating hyper-plane
- Sometimes $\vec{\phi}(\mathbf{x}_i)$ is an infinite dimensional vector so its good we don't have to compute it!

Kernel Trick

- If we choose a **positive semi-definite** kernel function $K(\mathbf{x}, \mathbf{y})$ then there exists functions $\phi_k(\mathbf{x})$, such that

$$K(\mathbf{x}_i, \mathbf{x}_j) = \vec{\phi}(\mathbf{x}_i) \cdot \vec{\phi}(\mathbf{x}_j)$$

(like an eigenvector decomposition of a matrix)

- Never need to compute $\phi_k(\mathbf{x}_i)$ explicitly as we only need the dot-product $\vec{\phi}(\mathbf{x}_i) \cdot \vec{\phi}(\mathbf{x}_j) = K(\mathbf{x}_i, \mathbf{x}_j)$ to compute maximum margin separating hyper-plane
- Sometimes $\vec{\phi}(\mathbf{x}_i)$ is an infinite dimensional vector so its good we don't have to compute it!

Kernel Functions

- Kernel functions are symmetric functions of two variable
- Strong restriction: *positive semi-definite*
- Examples

Quadratic kernel: $K(\mathbf{x}_1, \mathbf{x}_2) = (\mathbf{x}_1^\top \mathbf{x}_2)^2$

Gaussian (RBF) kernel: $K(\mathbf{x}_1, \mathbf{x}_2) = e^{-\gamma \|\mathbf{x}_1 - \mathbf{x}_2\|^2}$

- Consider the mapping

$$\mathbf{x}_i = \begin{pmatrix} x_i \\ y_i \end{pmatrix} \rightarrow \phi(\mathbf{x}_i) = \begin{pmatrix} x_i^2 \\ y_i^2 \\ \sqrt{2} x_i y_i \end{pmatrix}$$

Kernel Functions

- Kernel functions are symmetric functions of two variable
- Strong restriction: *positive semi-definite*
- Examples

Quadratic kernel: $K(\mathbf{x}_1, \mathbf{x}_2) = (\mathbf{x}_1^\top \mathbf{x}_2)^2$

Gaussian (RBF) kernel: $K(\mathbf{x}_1, \mathbf{x}_2) = e^{-\gamma \|\mathbf{x}_1 - \mathbf{x}_2\|^2}$

- Consider the mapping

$$\mathbf{x}_i = \begin{pmatrix} x_i \\ y_i \end{pmatrix} \rightarrow \phi(\mathbf{x}_i) = \begin{pmatrix} x_i^2 \\ y_i^2 \\ \sqrt{2} x_i y_i \end{pmatrix}$$

Kernel Functions

- Kernel functions are symmetric functions of two variable
- Strong restriction: *positive semi-definite*
- Examples

Quadratic kernel: $K(\mathbf{x}_1, \mathbf{x}_2) = (\mathbf{x}_1^\top \mathbf{x}_2)^2$

Gaussian (RBF) kernel: $K(\mathbf{x}_1, \mathbf{x}_2) = e^{-\gamma \|\mathbf{x}_1 - \mathbf{x}_2\|^2}$

- Consider the mapping

$$\mathbf{x}_i = \begin{pmatrix} x_i \\ y_i \end{pmatrix} \rightarrow \phi(\mathbf{x}_i) = \begin{pmatrix} x_i^2 \\ y_i^2 \\ \sqrt{2} x_i y_i \end{pmatrix}$$

Kernel Functions

- Kernel functions are symmetric functions of two variable
- Strong restriction: *positive semi-definite*
- Examples

Quadratic kernel: $K(\mathbf{x}_1, \mathbf{x}_2) = (\mathbf{x}_1^\top \mathbf{x}_2)^2$

Gaussian (RBF) kernel: $K(\mathbf{x}_1, \mathbf{x}_2) = e^{-\gamma \|\mathbf{x}_1 - \mathbf{x}_2\|^2}$

- Consider the mapping

$$\mathbf{x}_i = \begin{pmatrix} x_i \\ y_i \end{pmatrix} \rightarrow \phi(\mathbf{x}_i) = \begin{pmatrix} x_i^2 \\ y_i^2 \\ \sqrt{2} x_i y_i \end{pmatrix}$$

Non-linearly Separation of Data

$$K(\mathbf{x}_1, \mathbf{x}_2) = \begin{pmatrix} x_1^2 & y_1^2 & \sqrt{2} x_1 y_1 \end{pmatrix} \begin{pmatrix} x_2^2 \\ y_2^2 \\ \sqrt{2} x_2 y_2 \end{pmatrix}$$

Non-linearly Separation of Data

$$K(\mathbf{x}_1, \mathbf{x}_2) = \begin{pmatrix} x_1^2 & y_1^2 & \sqrt{2} x_1 y_1 \end{pmatrix} \begin{pmatrix} x_2^2 \\ y_2^2 \\ \sqrt{2} x_2 y_2 \end{pmatrix} = x_1^2 x_2^2 + y_1^2 y_2^2 + 2 x_1 y_1 x_2 y_2$$

Non-linearly Separation of Data


$$\begin{aligned} K(\mathbf{x}_1, \mathbf{x}_2) &= \begin{pmatrix} x_1^2 & y_1^2 & \sqrt{2} x_1 y_1 \end{pmatrix} \begin{pmatrix} x_2^2 \\ y_2^2 \\ \sqrt{2} x_2 y_2 \end{pmatrix} = x_1^2 x_2^2 + y_1^2 y_2^2 + 2 x_1 y_1 x_2 y_2 \\ &= (x_1 x_2 + y_1 y_2)^2 \end{aligned}$$

Non-linearly Separation of Data

$$\begin{aligned} K(\mathbf{x}_1, \mathbf{x}_2) &= \begin{pmatrix} x_1^2 & y_1^2 & \sqrt{2} x_1 y_1 \end{pmatrix} \begin{pmatrix} x_2^2 \\ y_2^2 \\ \sqrt{2} x_2 y_2 \end{pmatrix} = x_1^2 x_2^2 + y_1^2 y_2^2 + 2 x_1 y_1 x_2 y_2 \\ &= (x_1 x_2 + y_1 y_2)^2 = (\mathbf{x}_1^T \mathbf{x}_2)^2 \end{aligned}$$


Non-linearly Separable Data

$$K(\mathbf{x}_1, \mathbf{x}_2) = \begin{pmatrix} x_1^2 & y_1^2 & \sqrt{2} x_1 y_1 \end{pmatrix} \begin{pmatrix} x_2^2 \\ y_2^2 \\ \sqrt{2} x_2 y_2 \end{pmatrix} = x_1^2 x_2^2 + y_1^2 y_2^2 + 2 x_1 y_1 x_2 y_2$$
$$= (x_1 x_2 + y_1 y_2)^2 = (\mathbf{x}_1^T \mathbf{x}_2)^2$$



Non-linearly Separation of Data



$$K(\mathbf{x}_1, \mathbf{x}_2) = \begin{pmatrix} x_1^2 & y_1^2 & \sqrt{2} x_1 y_1 \end{pmatrix} \begin{pmatrix} x_2^2 \\ y_2^2 \\ \sqrt{2} x_2 y_2 \end{pmatrix} = x_1^2 x_2^2 + y_1^2 y_2^2 + 2 x_1 y_1 x_2 y_2$$
$$= (x_1 x_2 + y_1 y_2)^2 = (\mathbf{x}_1^T \mathbf{x}_2)^2$$



Non-linearly Separation of Data

$$K(\mathbf{x}_1, \mathbf{x}_2) = \begin{pmatrix} x_1^2 & y_1^2 & \sqrt{2} x_1 y_1 \end{pmatrix} \begin{pmatrix} x_2^2 \\ y_2^2 \\ \sqrt{2} x_2 y_2 \end{pmatrix} = x_1^2 x_2^2 + y_1^2 y_2^2 + 2 x_1 y_1 x_2 y_2$$



$$= (x_1 x_2 + y_1 y_2)^2 = (\mathbf{x}_1^T \mathbf{x}_2)^2$$



Non-linearly Separation of Data

$$K(\mathbf{x}_1, \mathbf{x}_2) = \begin{pmatrix} x_1^2 & y_1^2 & \sqrt{2} x_1 y_1 \end{pmatrix} \begin{pmatrix} x_2^2 \\ y_2^2 \\ \sqrt{2} x_2 y_2 \end{pmatrix} = x_1^2 x_2^2 + y_1^2 y_2^2 + 2 x_1 y_1 x_2 y_2$$



$$= (x_1 x_2 + y_1 y_2)^2 = (\mathbf{x}_1^T \mathbf{x}_2)^2$$



Non-linearly Separation of Data


$$K(\mathbf{x}_1, \mathbf{x}_2) = \begin{pmatrix} x_1^2 & y_1^2 & \sqrt{2} x_1 y_1 \end{pmatrix} \begin{pmatrix} x_2^2 \\ y_2^2 \\ \sqrt{2} x_2 y_2 \end{pmatrix} = x_1^2 x_2^2 + y_1^2 y_2^2 + 2 x_1 y_1 x_2 y_2$$

$$= (x_1 x_2 + y_1 y_2)^2 = (\mathbf{x}_1^T \mathbf{x}_2)^2$$



Outline

1. The Big Picture
2. **Practice**
3. Maximum Margins



Computing the Maximum-Margin Hyper-plane

- We will derive the formula for the minimum-margin hyper-plane in the next lecture
- This gives us a quadratic programming problem
- Through a neat trick we can represent this problem in a “dual form” where we
- Never need to compute $\phi_i(\mathbf{x})$ only need to compute $K(\mathbf{x}_i, \mathbf{x}_j)$
- When we use the kernel trick the time to compute the solution to the quadratic programming problem is $p N^3$ where N is the number of training examples and p is the number of features

Computing the Maximum-Margin Hyper-plane

- We will derive the formula for the minimum-margin hyper-plane in the next lecture
- This gives us a quadratic programming problem
- Through a neat trick we can represent this problem in a “dual form” where we
- Never need to compute $\phi_i(\mathbf{x})$ only need to compute $K(\mathbf{x}_i, \mathbf{x}_j)$
- When we use the kernel trick the time to compute the solution to the quadratic programming problem is $p N^3$ where N is the number of training examples and p is the number of features

Computing the Maximum-Margin Hyper-plane

- We will derive the formula for the minimum-margin hyper-plane in the next lecture
- This gives us a quadratic programming problem
- Through a neat trick we can represent this problem in a “dual form” where we
- Never need to compute $\phi_i(\mathbf{x})$ only need to compute $K(\mathbf{x}_i, \mathbf{x}_j)$
- When we use the kernel trick the time to compute the solution to the quadratic programming problem is $p N^3$ where N is the number of training examples and p is the number of features

Computing the Maximum-Margin Hyper-plane

- We will derive the formula for the minimum-margin hyper-plane in the next lecture
- This gives us a quadratic programming problem
- Through a neat trick we can represent this problem in a “dual form” where we
- Never need to compute $\phi_i(\mathbf{x})$ only need to compute $K(\mathbf{x}_i, \mathbf{x}_j)$
- When we use the kernel trick the time to compute the solution to the quadratic programming problem is $p N^3$ where N is the number of training examples and p is the number of features

Computing the Maximum-Margin Hyper-plane

- We will derive the formula for the minimum-margin hyper-plane in the next lecture
- This gives us a quadratic programming problem
- Through a neat trick we can represent this problem in a “dual form” where we
- Never need to compute $\phi_i(\mathbf{x})$ only need to compute $K(\mathbf{x}_i, \mathbf{x}_j)$
- When we use the kernel trick the time to compute the solution to the quadratic programming problem is $p N^3$ where N is the number of training examples and p is the number of features

Getting SVMs to Work Well


- SVMs rely on distances between data points
- These will change relative to each other if we rescale some features but not other—giving different maximum-margin hyper-planes

Getting SVMs to Work Well

- SVMs rely on distances between data points
- These will change relative to each other if we rescale some features but not other—giving different maximum-margin hyper-planes


Getting SVMs to Work Well

- SVMs rely on distances between data points
- These will change relative to each other if we rescale some features but not other—giving different maximum-margin hyper-planes




Getting SVMs to Work Well

- SVMs rely on distances between data points
- These will change relative to each other if we rescale some features but not other—giving different maximum-margin hyper-planes




Getting SVMs to Work Well

- SVMs rely on distances between data points
- These will change relative to each other if we rescale some features but not other—giving different maximum-margin hyper-planes




Getting SVMs to Work Well

- SVMs rely on distances between data points
- These will change relative to each other if we rescale some features but not other—giving different maximum-margin hyper-planes



Getting SVMs to Work Well

- SVMs rely on distances between data points
- These will change relative to each other if we rescale some features but not other—giving different maximum-margin hyper-planes




- If we don't know what features are important (most often the case), then it is worth scaling each feature (for example, so their range is between 0 and 1 or their variance is 1)

Soft Margins

- Sometimes the margin constraint is too severe
- Relax constraints by introducing *slack variables*, $\xi_k \geq 0$

$$y_k(\mathbf{x}_k^\top \mathbf{w} - b) \geq 1 - \xi_k$$




- Minimise $\frac{\|\mathbf{w}\|^2}{2} + C \sum_{k=1}^n \xi_k$ subject to constraints
- Large C punishes slack variables

Soft Margins

- Sometimes the margin constraint is too severe
- Relax constraints by introducing *slack variables*, $\xi_k \geq 0$

$$y_k(\mathbf{x}_k^\top \mathbf{w} - b) \geq 1 - \xi_k$$

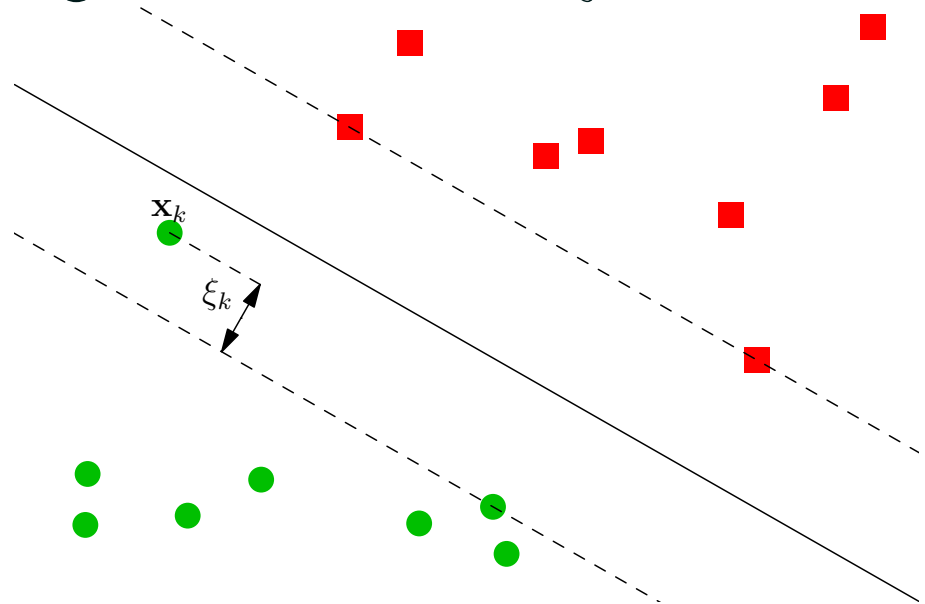


- Minimise $\frac{\|\mathbf{w}\|^2}{2} + C \sum_{k=1}^n \xi_k$ subject to constraints
- Large C punishes slack variables

Soft Margins

- Sometimes the margin constraint is too severe
- Relax constraints by introducing *slack variables*, $\xi_k \geq 0$

$$y_k(\mathbf{x}_k^\top \mathbf{w} - b) \geq 1 - \xi_k$$




- Minimise $\frac{\|\mathbf{w}\|^2}{2} + C \sum_{k=1}^n \xi_k$ subject to constraints
- Large C punishes slack variables

Soft Margins

- Sometimes the margin constraint is too severe
- Relax constraints by introducing *slack variables*, $\xi_k \geq 0$

$$y_k(\mathbf{x}_k^\top \mathbf{w} - b) \geq 1 - \xi_k$$



- Minimise $\frac{\|\mathbf{w}\|^2}{2} + C \sum_{k=1}^n \xi_k$ subject to constraints
- Large C punishes slack variables

Optimising C

- In practice it can make a huge difference to the performance if we change C
- Optimal C values changes by many orders of magnitude e.g. 2^{-5} – 2^{15}
- Typically optimised by a grid search (start from 2^{-5} say and double until you reach 2^{15})

Optimising C

- In practice it can make a huge difference to the performance if we change C
- Optimal C values changes by many orders of magnitude e.g. 2^{-5} – 2^{15}
- Typically optimised by a grid search (start from 2^{-5} say and double until you reach 2^{15})

Optimising C

- In practice it can make a huge difference to the performance if we change C
- Optimal C values changes by many orders of magnitude e.g. 2^{-5} – 2^{15}
- Typically optimised by a grid search (start from 2^{-5} say and double until you reach 2^{15})

Choosing the Right Kernel Function

- There are kernels design for particular data types (e.g. string kernels for text or biological sequences)
- For numerical data people tend to look at using no kernel (linear SVM), a radial basis function (Gaussian) kernel or polynomial kernels
- Kernel's often come with parameters, e.g. the popular radial basis function kernel

$$K(\mathbf{x}, \mathbf{y}) = e^{-\gamma \|\mathbf{x} - \mathbf{y}\|^2}$$

- Optimal γ values range over 2^{-15} – 2^3

Choosing the Right Kernel Function

- There are kernels design for particular data types (e.g. string kernels for text or biological sequences)
- For numerical data people tend to look at using no kernel (linear SVM), a radial basis function (Gaussian) kernel or polynomial kernels
- Kernel's often come with parameters, e.g. the popular radial basis function kernel

$$K(\mathbf{x}, \mathbf{y}) = e^{-\gamma \|\mathbf{x} - \mathbf{y}\|^2}$$

- Optimal γ values range over 2^{-15} – 2^3

Choosing the Right Kernel Function

- There are kernels design for particular data types (e.g. string kernels for text or biological sequences)
- For numerical data people tend to look at using no kernel (linear SVM), a radial basis function (Gaussian) kernel or polynomial kernels
- Kernel's often come with parameters, e.g. the popular radial basis function kernel

$$K(\mathbf{x}, \mathbf{y}) = e^{-\gamma \|\mathbf{x} - \mathbf{y}\|^2}$$

- Optimal γ values range over 2^{-15} – 2^3

Choosing the Right Kernel Function

- There are kernels design for particular data types (e.g. string kernels for text or biological sequences)
- For numerical data people tend to look at using no kernel (linear SVM), a radial basis function (Gaussian) kernel or polynomial kernels
- Kernel's often come with parameters, e.g. the popular radial basis function kernel

$$K(\mathbf{x}, \mathbf{y}) = e^{-\gamma \|\mathbf{x} - \mathbf{y}\|^2}$$

- Optimal γ values range over 2^{-15} – 2^3

SVM Libraries

- Although SVMs have unique solutions, they require very well written optimisers
- If you have a large data set they can be very slow
- There are good libraries out there, svmlib, svm-lite, etc.
- These will often automate normalisation of data and grid search for parameters

SVM Libraries

- Although SVMs have unique solutions, they require very well written optimisers
- If you have a large data set they can be very slow
- There are good libraries out there, svmlib, svm-lite, etc.
- These will often automate normalisation of data and grid search for parameters

SVM Libraries


- Although SVMs have unique solutions, they require very well written optimisers
- If you have a large data set they can be very slow
- There are good libraries out there, svmlib, svm-lite, etc.
- These will often automate normalisation of data and grid search for parameters

SVM Libraries

- Although SVMs have unique solutions, they require very well written optimisers
- If you have a large data set they can be very slow
- There are good libraries out there, svm-lib, svm-lite, etc.
- These will often automate normalisation of data and grid search for parameters

Outline

1. The Big Picture
2. Practice
3. **Maximum Margins**



Dot Product

- Recall the dot product

$$\left(\text{---} \right) \left(\begin{array}{c} | \\ | \\ | \end{array} \right) = \left(\blacksquare \right)$$

$$x \cdot y = x^{\top} y$$

Dot Product

- Recall the dot product

$$\left(\text{---} \right) \left(\begin{array}{c} | \\ | \\ | \end{array} \right) = \left(\blacksquare \right)$$

$$\mathbf{x} \cdot \mathbf{y} = \mathbf{x}^\top \mathbf{y} = \sum_{i=1}^n x_i y_i$$

Dot Product

- Recall the dot product

$$\left(\text{---} \right) \left(\begin{array}{c} | \\ | \\ | \end{array} \right) = \left(\blacksquare \right)$$

$$\mathbf{x} \cdot \mathbf{y} = \mathbf{x}^\top \mathbf{y} = \sum_{i=1}^n x_i y_i = \|\mathbf{x}\| \|\mathbf{y}\| \cos(\theta)$$

Dot Product

- Recall the dot product

$$\left(\text{---} \right) \left(\begin{pmatrix} | \\ | \\ | \end{pmatrix} \right) = \left(\blacksquare \right) \quad \mathbf{x} \cdot \mathbf{y} = \mathbf{x}^\top \mathbf{y} = \sum_{i=1}^n x_i y_i = \|\mathbf{x}\| \|\mathbf{y}\| \cos(\theta)$$


- If $\|\mathbf{w}\| = 1$ then $\mathbf{x}^\top \mathbf{w} = \|\mathbf{x}\| \cos(\theta)$

Dot Product

- Recall the dot product

$$\left(\text{---} \right) \left(\begin{pmatrix} | \\ | \\ | \end{pmatrix} \right) = \left(\begin{pmatrix} \blacksquare \end{pmatrix} \right) \quad x \cdot y = x^\top y = \sum_{i=1}^n x_i y_i = \|x\| \|y\| \cos(\theta)$$

- If $\|w\| = 1$ then $x^\top w = \|x\| \cos(\theta)$




Dot Product

- Recall the dot product

$$\left(\text{---} \right) \left(\begin{pmatrix} | \\ | \\ | \end{pmatrix} \right) = \left(\blacksquare \right) \quad x \cdot y = x^\top y = \sum_{i=1}^n x_i y_i = \|x\| \|y\| \cos(\theta)$$

- If $\|w\| = 1$ then $x^\top w = \|x\| \cos(\theta)$




Dot Product

- Recall the dot product

$$\left(\text{---} \right) \left(\begin{pmatrix} | \\ | \\ | \end{pmatrix} \right) = \left(\blacksquare \right) \quad x \cdot y = x^\top y = \sum_{i=1}^n x_i y_i = \|x\| \|y\| \cos(\theta)$$

- If $\|w\| = 1$ then $x^\top w = \|x\| \cos(\theta)$




Dot Product

- Recall the dot product

$$\left(\text{---} \right) \left(\begin{pmatrix} | \\ | \\ | \end{pmatrix} \right) = \left(\blacksquare \right) \quad x \cdot y = x^\top y = \sum_{i=1}^n x_i y_i = \|x\| \|y\| \cos(\theta)$$

- If $\|w\| = 1$ then $x^\top w = \|x\| \cos(\theta)$




Dot Product

- Recall the dot product

$$\left(\text{---} \right) \left(\begin{pmatrix} | \\ | \\ | \end{pmatrix} \right) = \left(\blacksquare \right) \quad x \cdot y = x^\top y = \sum_{i=1}^n x_i y_i = \|x\| \|y\| \cos(\theta)$$

- If $\|w\| = 1$ then $x^\top w = \|x\| \cos(\theta)$




Dot Product

- Recall the dot product

$$\left(\text{---} \right) \left(\begin{pmatrix} | \\ | \\ | \end{pmatrix} \right) = \left(\blacksquare \right) \quad x \cdot y = x^\top y = \sum_{i=1}^n x_i y_i = \|x\| \|y\| \cos(\theta)$$

- If $\|w\| = 1$ then $x^\top w = \|x\| \cos(\theta)$




Dot Product

- Recall the dot product

$$\left(\text{---} \right) \left(\begin{pmatrix} | \\ | \\ | \end{pmatrix} \right) = \left(\blacksquare \right) \quad x \cdot y = x^\top y = \sum_{i=1}^n x_i y_i = \|x\| \|y\| \cos(\theta)$$

- If $\|w\| = 1$ then $x^\top w = \|x\| \cos(\theta)$




Dot Product

- Recall the dot product

$$\left(\text{---} \right) \left(\begin{pmatrix} | \\ | \\ | \end{pmatrix} \right) = \left(\blacksquare \right) \quad x \cdot y = x^T y = \sum_{i=1}^n x_i y_i = \|x\| \|y\| \cos(\theta)$$

- If $\|w\| = 1$ then $x^T w = \|x\| \cos(\theta)$



Maximise Margin

- Consider a linearly separable set of data
 - ★ $\mathcal{D} = \{(\mathbf{x}_k, y_k)\}_{k=1}^P$
 - ★ $y_k \in \{-1, 1\}$
- Our task is to find a separating plane defined by the orthogonal vector \mathbf{w} and a threshold b such that

$$y_k \left(\frac{\mathbf{w}^\top \mathbf{x}_k}{\|\mathbf{w}\|} - b \right) \geq m$$

where m is the margin


Maximise Margin

- Consider a linearly separable set of data
 - ★ $\mathcal{D} = \{(\mathbf{x}_k, y_k)\}_{k=1}^P$
 - ★ $y_k \in \{-1, 1\}$
- Our task is to find a separating plane defined by the orthogonal vector \mathbf{w} and a threshold b such that


$$y_k \left(\frac{\mathbf{w}^\top \mathbf{x}_k}{\|\mathbf{w}\|} - b \right) \geq m$$

where m is the margin


Distance to hyperplanes




Distance to hyperplanes




Distance to hyperplanes




Distance to hyperplanes




Distance to hyperplanes




Distance to hyperplanes




Distance to hyperplanes



Distance to hyperplanes



Distance to hyperplanes



Constrained Optimisation

- Wish to find \mathbf{w} and b to maximise m subject to constraints

$$y_k \left(\frac{\mathbf{w}^\top \mathbf{x}_k}{\|\mathbf{w}\|} - b \right) \geq m \quad \text{for all } k = 1, 2, \dots, P$$

- If we divide through by m

$$y_k \left(\frac{\mathbf{w}^\top \mathbf{x}_k}{m \|\mathbf{w}\|} - \frac{b}{m} \right) \geq 1 \quad \text{for all } k = 1, 2, \dots, P$$

- Define $\hat{\mathbf{w}} = \mathbf{w}/(m\|\mathbf{w}\|)$ and $\hat{b} = b/m$

$$y_k \left(\hat{\mathbf{w}}^\top \mathbf{x}_k - \hat{b} \right) \geq 1$$

Constrained Optimisation

- Wish to find \mathbf{w} and b to maximise m subject to constraints

$$y_k \left(\frac{\mathbf{w}^\top \mathbf{x}_k}{\|\mathbf{w}\|} - b \right) \geq m \quad \text{for all } k = 1, 2, \dots, P$$

- If we divide through by m

$$y_k \left(\frac{\mathbf{w}^\top \mathbf{x}_k}{m \|\mathbf{w}\|} - \frac{b}{m} \right) \geq 1 \quad \text{for all } k = 1, 2, \dots, P$$

- Define $\hat{\mathbf{w}} = \mathbf{w}/(m\|\mathbf{w}\|)$ and $\hat{b} = b/m$

$$y_k \left(\hat{\mathbf{w}}^\top \mathbf{x}_k - \hat{b} \right) \geq 1$$

Constrained Optimisation

- Wish to find \mathbf{w} and b to maximise m subject to constraints

$$y_k \left(\frac{\mathbf{w}^\top \mathbf{x}_k}{\|\mathbf{w}\|} - b \right) \geq m \quad \text{for all } k = 1, 2, \dots, P$$

- If we divide through by m

$$y_k \left(\frac{\mathbf{w}^\top \mathbf{x}_k}{m \|\mathbf{w}\|} - \frac{b}{m} \right) \geq 1 \quad \text{for all } k = 1, 2, \dots, P$$

- Define $\hat{\mathbf{w}} = \mathbf{w}/(m\|\mathbf{w}\|)$ and $\hat{b} = b/m$

$$y_k \left(\hat{\mathbf{w}}^\top \mathbf{x}_k - \hat{b} \right) \geq 1$$

Quadratic Programming Problem

- Note that as $\hat{\mathbf{w}} = \mathbf{w} / (m \|\mathbf{w}\|)$

$$\|\hat{\mathbf{w}}\| = \left\| \frac{\mathbf{w}}{m \|\mathbf{w}\|} \right\| = \frac{1}{m}$$

- Minimising $\|\hat{\mathbf{w}}\|^2$ is equivalent to maximising the margin m
- Can write the optimisation problem as a *quadratic programming problem*

$$\min_{\hat{\mathbf{w}}, \hat{b}} \frac{\|\hat{\mathbf{w}}\|^2}{2} \quad \text{subject to } y_k \left(\hat{\mathbf{w}}^\top \mathbf{x}_k - \hat{b} \right) \geq 1 \text{ for all } k = 1, 2, \dots, P$$

Quadratic Programming Problem

- Note that as $\hat{\mathbf{w}} = \mathbf{w}/(m\|\mathbf{w}\|)$

$$\|\hat{\mathbf{w}}\| = \left\| \frac{\mathbf{w}}{m\|\mathbf{w}\|} \right\| = \frac{1}{m}$$

- Minimising $\|\hat{\mathbf{w}}\|^2$ is equivalent to maximising the margin m
- Can write the optimisation problem as a *quadratic programming problem*

$$\min_{\hat{\mathbf{w}}, \hat{b}} \frac{\|\hat{\mathbf{w}}\|^2}{2} \quad \text{subject to } y_k \left(\hat{\mathbf{w}}^\top \mathbf{x}_k - \hat{b} \right) \geq 1 \text{ for all } k = 1, 2, \dots, P$$

Quadratic Programming Problem

- Note that as $\hat{\mathbf{w}} = \mathbf{w}/(m\|\mathbf{w}\|)$

$$\|\hat{\mathbf{w}}\| = \left\| \frac{\mathbf{w}}{m\|\mathbf{w}\|} \right\| = \frac{1}{m}$$


- Minimising $\|\hat{\mathbf{w}}\|^2$ is equivalent to maximising the margin m
- Can write the optimisation problem as a *quadratic programming problem*

$$\min_{\hat{\mathbf{w}}, \hat{b}} \frac{\|\hat{\mathbf{w}}\|^2}{2} \quad \text{subject to } y_k \left(\hat{\mathbf{w}}^\top \mathbf{x}_k - \hat{b} \right) \geq 1 \text{ for all } k = 1, 2, \dots, P$$

Convexity

- The quadratic function $f(x) = x^2$ is an example of a *convex* function satisfying

$$f(tx + (1 - t)y) \leq tf(x) + (1 - t)f(y) \quad \text{for } 0 \leq t \leq 1$$




- This extends to high dimensions

$$f(t\mathbf{x} + (1 - t)\mathbf{y}) \leq tf(\mathbf{x}) + (1 - t)f(\mathbf{y})$$

Convexity

- The quadratic function $f(x) = x^2$ is an example of a *convex* function satisfying

$$f(tx + (1 - t)y) \leq tf(x) + (1 - t)f(y) \quad \text{for } 0 \leq t \leq 1$$




- This extends to high dimensions

$$f(t\mathbf{x} + (1 - t)\mathbf{y}) \leq tf(\mathbf{x}) + (1 - t)f(\mathbf{y})$$


Unique Minimum

- Convex function have a unique minimum
- The existence of a local minimum would break convexity
 - ★ The line connecting a local minimum to a global minimum would be strictly decreasing
 - ★ Thus there are points next to the local minimum with lower values
 - ★ This is a contradiction
- This remains true if we consider convex functions that a constrained to live in a **convex region**




Unique Minimum

- Convex function have a unique minimum
- The existence of a local minimum would break convexity
 - ★ The line connecting a local minimum to a global minimum would be strictly decreasing
 - ★ Thus there are points next to the local minimum with lower values
 - ★ This is a contradiction
- This remains true if we consider convex functions that a constrained to live in a **convex region**




Unique Minimum

- Convex function have a unique minimum
- The existence of a local minimum would break convexity
 - ★ The line connecting a local minimum to a global minimum would be strictly decreasing
 - ★ Thus there are points next to the local minimum with lower values
 - ★ This is a contradiction
- This remains true if we consider convex functions that a constrained to live in a **convex region**




Unique Minimum

- Convex function have a unique minimum
- The existence of a local minimum would break convexity
 - ★ The line connecting a local minimum to a global minimum would be strictly decreasing
 - ★ Thus there are points next to the local minimum with lower values
 - ★ This is a contradiction
- This remains true if we consider convex functions that a constrained to live in a **convex region**




Unique Minimum

- Convex function have a unique minimum
- The existence of a local minimum would break convexity
 - ★ The line connecting a local minimum to a global minimum would be strictly decreasing
 - ★ Thus there are points next to the local minimum with lower values
 - ★ This is a contradiction
- This remains true if we consider convex functions that a constrained to live in a **convex region**




Unique Minimum

- Convex function have a unique minimum
- The existence of a local minimum would break convexity
 - ★ The line connecting a local minimum to a global minimum would be strictly decreasing
 - ★ Thus there are points next to the local minimum with lower values
 - ★ This is a contradiction
- This remains true if we consider convex functions that a constrained to live in a **convex region**



Convex Regions


- Convex regions are familiar



- Convex function constrained to lie in a convex region will also have unique minima as minima can't hide in corners unlike concave regions
- Quadratic programming problems involving a quadratic function and linear constraints are convex and have a unique minimum

Convex Regions


- Convex regions are familiar



- Convex function constrained to lie in a convex region will also have unique minima as minima can't hide in corners unlike concave regions
- Quadratic programming problems involving a quadratic function and linear constraints are convex and have a unique minimum


Convex Regions

- Convex regions are familiar




- Convex function constrained to lie in a convex region will also have unique minima as minima can't hide in corners unlike concave regions
- Quadratic programming problems involving a quadratic function and linear constraints are convex and have a unique minimum


Quadratic Programming in SVMs




Quadratic Programming in SVMs




Quadratic Programming in SVMs



Quadratic Programming in SVMs



Quadratic Programming in SVMs



Quadratic Programming

- We have a quadratic programming problem for the weights $\hat{\mathbf{w}} = (\hat{w}_1, \hat{w}_2, \dots, \hat{w}_n)$ and bias b and P constraints
- This is a classic but fiddly optimisation problems
- It can be solved in $O(n^3)$ time (it involves inverting matrices)
- We will see that there is an equivalent dual problem which allows us to use the kernel trick

Quadratic Programming

- We have a quadratic programming problem for the weights $\hat{\mathbf{w}} = (\hat{w}_1, \hat{w}_2, \dots, \hat{w}_n)$ and bias b and P constraints
- This is a classic but fiddly optimisation problems
- It can be solved in $O(n^3)$ time (it involves inverting matrices)
- We will see that there is an equivalent dual problem which allows us to use the kernel trick

Quadratic Programming

- We have a quadratic programming problem for the weights $\hat{\mathbf{w}} = (\hat{w}_1, \hat{w}_2, \dots, \hat{w}_n)$ and bias b and P constraints
- This is a classic but fiddly optimisation problems
- It can be solved in $O(n^3)$ time (it involves inverting matrices)
- We will see that there is an equivalent dual problem which allows us to use the kernel trick

Quadratic Programming

- We have a quadratic programming problem for the weights $\hat{\mathbf{w}} = (\hat{w}_1, \hat{w}_2, \dots, \hat{w}_n)$ and bias b and P constraints
- This is a classic but fiddly optimisation problems
- It can be solved in $O(n^3)$ time (it involves inverting matrices)
(pew it is not NP-complete!)
- We will see that there is an equivalent dual problem which allows us to use the kernel trick


Quadratic Programming

- We have a quadratic programming problem for the weights $\hat{\mathbf{w}} = (\hat{w}_1, \hat{w}_2, \dots, \hat{w}_n)$ and bias b and P constraints
- This is a classic but fiddly optimisation problems
- It can be solved in $O(n^3)$ time (it involves inverting matrices) (phew it is not NP-complete!)
- We will see that there is an equivalent dual problem which allows us to use the kernel trick

Soft Margins

- Can relax constraints by introducing *slack variables*, $\xi_k \geq 0$

$$y_k(\hat{\mathbf{w}}^\top \mathbf{x}_k - \hat{b}) \geq 1 - \xi_k$$




- Minimise $\frac{\|\hat{\mathbf{w}}\|^2}{2} + C \sum_{k=1}^n \xi_k$ subject to constraints
- We've added a linear function that leaves our objective convex

Soft Margins

- Can relax constraints by introducing *slack variables*, $\xi_k \geq 0$

$$y_k(\hat{\mathbf{w}}^\top \mathbf{x}_k - \hat{b}) \geq 1 - \xi_k$$




- Minimise $\frac{\|\hat{\mathbf{w}}\|^2}{2} + C \sum_{k=1}^n \xi_k$ subject to constraints
- We've added a linear function that leaves our objective convex

Soft Margins

- Can relax constraints by introducing *slack variables*, $\xi_k \geq 0$

$$y_k(\hat{\mathbf{w}}^\top \mathbf{x}_k - \hat{b}) \geq 1 - \xi_k$$



- Minimise $\frac{\|\hat{\mathbf{w}}\|^2}{2} + C \sum_{k=1}^n \xi_k$ subject to constraints
- We've added a linear function that leaves our objective convex

Convexity of Lasso

- Recall that in Lasso we are asked to minimise

$$E = \sum_{n=1}^N (\mathbf{w}^\top \mathbf{x}_n - y_n)^2 + \nu \sum_{i=1}^p |w_i|$$

- We can rewrite this (using $w_i = w_i^+ - w_i^-$) as

$$E = \sum_{n=1}^N ((\mathbf{w}^+ - \mathbf{w}^-)^\top \mathbf{x}_n - y_n)^2 + \nu \sum_{i=1}^p (w_i^+ + w_i^-)$$

subject to $w_i^+, w_i^- \geq 0$ for all i

- Again it Lasso has a unique minima

Convexity of Lasso

- Recall that in Lasso we are asked to minimise

$$E = \sum_{n=1}^N (\mathbf{w}^\top \mathbf{x}_n - y_n)^2 + \nu \sum_{i=1}^p |w_i|$$

- We can rewrite this (using $w_i = w_i^+ - w_i^-$) as

$$E = \sum_{n=1}^N ((\mathbf{w}^+ - \mathbf{w}^-)^\top \mathbf{x}_n - y_n)^2 + \nu \sum_{i=1}^p (w_i^+ + w_i^-)$$

subject to $w_i^+, w_i^- \geq 0$ for all i

- Again it Lasso has a unique minima

Convexity of Lasso

- Recall that in Lasso we are asked to minimise

$$E = \sum_{n=1}^N (\mathbf{w}^\top \mathbf{x}_n - y_n)^2 + \nu \sum_{i=1}^p |w_i|$$

- We can rewrite this (using $w_i = w_i^+ - w_i^-$) as

$$E = \sum_{n=1}^N ((\mathbf{w}^+ - \mathbf{w}^-)^\top \mathbf{x}_n - y_n)^2 + \nu \sum_{i=1}^p (w_i^+ + w_i^-)$$

subject to $w_i^+, w_i^- \geq 0$ for all i

- Again it Lasso has a unique minima

Conclusions

- We've seen how SVMs work
- We've learnt how to use them
- We've seen that we can find the maximum margin hyper-plane by solving a quadratic programming problem (with a unique solution)
- This is a convex optimisation problem with a unique optimum
- Next we will look at the **dual problem**

Conclusions

- We've seen how SVMs work
- We've learnt how to use them
- We've seen that we can find the maximum margin hyper-plane by solving a quadratic programming problem (with a unique solution)
- This is a convex optimisation problem with a unique optimum
- Next we will look at the **dual problem**

Conclusions

- We've seen how SVMs work
- We've learnt how to use them
- We've seen that we can find the maximum margin hyper-plane by solving a quadratic programming problem (with a unique solution)
- This is a convex optimisation problem with a unique optimum
- Next we will look at the **dual problem**

Conclusions

- We've seen how SVMs work
- We've learnt how to use them
- We've seen that we can find the maximum margin hyper-plane by solving a quadratic programming problem (with a unique solution)
- This is a convex optimisation problem with a unique optimum
- Next we will look at the **dual problem**

Conclusions

- We've seen how SVMs work
- We've learnt how to use them
- We've seen that we can find the maximum margin hyper-plane by solving a quadratic programming problem (with a unique solution)
- This is a convex optimisation problem with a unique optimum
- Next we will look at the **dual problem**