

Support Vector Machines, maximum margins

Support Vector Machines

- Support vector machines, when used right, often have the best generalisation results
- They are typically used on numerical data, but can and have been adapted to text, sequences, etc.
- Although not as trendy as deep learning, they will often be the method of choice on small data sets
- They subtly regularise themselves, choosing a solution that generalises well from a host of different problems

Extended Feature Space

- To increase the likelihood of linear-separability we often use a high-dimensional mapping

$$\mathbf{x} = (x_1, x_2, \dots, x_p) \rightarrow \vec{\phi}(\mathbf{x}) = (\phi_1(\mathbf{x}), \phi_2(\mathbf{x}), \dots, \phi_r(\mathbf{x}))$$

$$r \gg p$$

- Finding the maximum margin hyper-plane is time consuming in “primal” form if r is large
- We can work in the “dual” space of patterns, then we only need to compute dot products

$$\vec{\phi}(\mathbf{x}_i) \cdot \vec{\phi}(\mathbf{x}_j) = \sum_{k=1}^r \phi_k(\mathbf{x}_i) \phi_k(\mathbf{x}_j)$$

Kernel Functions

- Kernel functions are symmetric functions of two variable
- Strong restriction: *positive semi-definite*
- Examples

Quadratic kernel: $K(\mathbf{x}_1, \mathbf{x}_2) = (\mathbf{x}_1^T \mathbf{x}_2)^2$

Gaussian (RBF) kernel: $K(\mathbf{x}_1, \mathbf{x}_2) = e^{-\gamma \|\mathbf{x}_1 - \mathbf{x}_2\|^2}$


- Consider the mapping

$$\mathbf{x}_i = \begin{pmatrix} x_i \\ y_i \end{pmatrix} \rightarrow \phi(\mathbf{x}_i) = \begin{pmatrix} x_i^2 \\ y_i^2 \\ \sqrt{2} x_i y_i \end{pmatrix}$$

1. The Big Picture


2. Practice

3. Maximum Margins



Linear Separation of Data

- SVMs classify linearly separable data



- Finds maximum-margin separating plane

Kernel Trick

- If we choose a **positive semi-definite** kernel function $K(\mathbf{x}, \mathbf{y})$ then there exists functions $\phi_k(\mathbf{x})$, such that

$$K(\mathbf{x}_i, \mathbf{x}_j) = \vec{\phi}(\mathbf{x}_i) \cdot \vec{\phi}(\mathbf{x}_j)$$



(like an eigenvector decomposition of a matrix)

- Never need to compute $\phi_k(\mathbf{x}_i)$ explicitly as we only need the dot-product $\vec{\phi}(\mathbf{x}_i) \cdot \vec{\phi}(\mathbf{x}_j) = K(\mathbf{x}_i, \mathbf{x}_j)$ to compute maximum margin separating hyper-plane
- Sometimes $\vec{\phi}(\mathbf{x}_i)$ is an infinite dimensional vector so its good we don't have to compute it!


Non-linearly Separation of Data

$$K(\mathbf{x}_1, \mathbf{x}_2) = \begin{pmatrix} x_1^2 & y_1^2 & \sqrt{2} x_1 y_1 \end{pmatrix} \begin{pmatrix} x_2^2 \\ y_2^2 \\ \sqrt{2} x_2 y_2 \end{pmatrix} = x_1^2 x_2^2 + y_1^2 y_2^2 + 2 x_1 y_1 x_2 y_2$$

$$= (x_1 x_2 + y_1 y_2)^2 = (\mathbf{x}_1^T \mathbf{x}_2)^2$$




1. The Big Picture
2. Practice
3. Maximum Margins



Getting SVMs to Work Well

- SVMs rely on distances between data points
- These will change relative to each other if we rescale some features but not other—giving different maximum-margin hyper-planes



- If we don't know what features are important (most often the case), then it is worth scaling each feature (for example, so their range is between 0 and 1 or their variance is 1)

Optimising C

- In practice it can make a huge difference to the performance if we change C
- Optimal C values changes by many orders of magnitude e.g. 2^{-5} – 2^{15}
- Typically optimised by a grid search (start from 2^{-5} say and double until you reach 2^{15})

SVM Libraries


- Although SVMs have unique solutions, they require very well written optimisers
- If you have a large data set they can be very slow
- There are good libraries out there, svm-lib, svm-lite, etc.
- These will often automate normalisation of data and grid search for parameters

- We will derive the formula for the minimum-margin hyper-plane in the next lecture
- This gives us a quadratic programming problem
- Through a neat trick we can represent this problem in a “dual form” where we
- Never need to compute $\phi_i(x)$ only need to compute $K(x_i, x_j)$
- When we use the kernel trick the time to compute the solution to the quadratic programming problem is $p m^3$ where m is the number of training examples and p is the number of features

Soft Margins

- Sometimes the margin constraint is too severe
- Relax constraints by introducing *slack variables*, $\xi_k \geq 0$

$$y_k(x_k^T w - b) \geq 1 - \xi_k$$



- Minimise $\frac{\|w\|^2}{2} + C \sum_{k=1}^n \xi_k$ subject to constraints
- Large C punishes slack variables

Choosing the Right Kernel Function


- There are kernels design for particular data types (e.g. string kernels for text or biological sequences)
- For numerical data people tend to look at using no kernel (linear SVM), a radial basis function (Gaussian) kernel or polynomial kernels
- Kernel's often come with parameters, e.g. the popular radial basis function kernel

$$K(x, y) = e^{-\gamma \|x - y\|^2}$$

- Optimal γ values range over 2^{-15} – 2^3

Outline

1. The Big Picture
2. Practice
3. Maximum Margins




Dot Product


- Recall the dot product

$$\left(\begin{array}{c} \text{---} \\ \text{---} \end{array} \right) \left(\begin{array}{c} \text{---} \\ \text{---} \end{array} \right) = \text{---} \quad \mathbf{x} \cdot \mathbf{y} = \mathbf{x}^T \mathbf{y} = \sum_{i=1}^n x_i y_i = \|\mathbf{x}\| \|\mathbf{y}\| \cos(\theta)$$

- If $\|\mathbf{w}\| = 1$ then $\mathbf{x}^T \mathbf{w} = \|\mathbf{x}\| \cos(\theta)$



Distance to hyperplanes



Quadratic Programming Problem

- Note that as $\hat{\mathbf{w}} = \mathbf{w} / (\gamma_m \|\mathbf{w}\|)$

$$\|\hat{\mathbf{w}}\| = \left\| \frac{\mathbf{w}}{\gamma_m \|\mathbf{w}\|} \right\| = \frac{1}{\gamma_m}$$

- Minimising $\|\hat{\mathbf{w}}\|^2$ is equivalent to maximising the margin γ_m
- Can write the optimisation problem as a *quadratic programming problem*

$$\min_{\hat{\mathbf{w}}, \hat{b}} \frac{\|\hat{\mathbf{w}}\|^2}{2} \quad \text{subject to } y_k (\hat{\mathbf{w}}^T \mathbf{x}_k - \hat{b}) \geq 1 \text{ for all } k = 1, 2, \dots, P$$


Unique Minimum

- Convex function have a unique minimum
- The existence of a local minimum would break convexity
 - The line connecting a local minimum to a global minimum would be strictly decreasing

- Thus there are points next to the local minimum with lower values

- This is a contradiction

- This remains true if we consider convex functions that are constrained to live in a **convex region**



Maximise Margin

- Consider a linearly separable set of data

$$\star \mathcal{D} = \{(\mathbf{x}_k, y_k)\}_{k=1}^P$$

$$\star y_k \in \{-1, 1\}$$

- Our task is to find a separating plane defined by the orthogonal vector \mathbf{w} and a threshold b such that

$$y_k \left(\frac{\mathbf{w}^T \mathbf{x}_k}{\|\mathbf{w}\|} - b \right) \geq \gamma_m$$

where γ_m is the margin

Constrained Optimisation

- Wish to find \mathbf{w} and b to maximise γ_m subject to constraints

$$y_k \left(\frac{\mathbf{w}^T \mathbf{x}_k}{\|\mathbf{w}\|} - b \right) \geq \gamma_m \quad \text{for all } k = 1, 2, \dots, P$$

- If we divide through by γ_m

$$y_k \left(\frac{\mathbf{w}^T \mathbf{x}_k}{\gamma_m \|\mathbf{w}\|} - \frac{b}{\gamma_m} \right) \geq 1 \quad \text{for all } k = 1, 2, \dots, P$$


- Define $\hat{\mathbf{w}} = \mathbf{w} / (\gamma_m \|\mathbf{w}\|)$ and $\hat{b} = b / \gamma_m$

$$y_k (\hat{\mathbf{w}}^T \mathbf{x}_k - \hat{b}) \geq 1$$

Convexity

- The quadratic function $f(x) = x^2$ is an example of a *convex* function satisfying


$$f(tx + (1-t)y) \leq tf(x) + (1-t)f(y) \quad \text{for } 0 \leq t \leq 1$$



- This extends to high dimensions
- $$f(tx + (1-t)y) \leq tf(\mathbf{x}) + (1-t)f(\mathbf{y})$$


Convex Regions

- Convex regions are familiar



- Convex function constrained to lie in a convex region will also have unique minima as minima can't hide in corners unlike concave regions


- Quadratic programming problems involving a quadratic function and linear constraints are convex and have a unique minimum



Soft Margins

- Can relax constraints by introducing *slack variables*, $\xi_k \geq 0$

$$y_k(\mathbf{w}^\top \mathbf{x}_k - \hat{b}) \geq 1 - \xi_k$$



- Minimise $\frac{\|\mathbf{w}\|^2}{2} + C \sum_{k=1}^m \xi_k$ subject to constraints
- We've added a linear function that leaves our objective convex

Conclusions

- We've seen how SVMs work
- We've learnt how to use them
- We've seen that we can find the maximum margin hyper-plane by solving a quadratic programming problem (with a unique solution)
- This is a convex optimisation problem with a unique optimum
- Next we will look at the **dual problem**

- We have a quadratic programming problem for the weights $\mathbf{w} = (w_1, w_2, \dots, w_n)$ and bias \hat{b} and m constraints
- This is a classic but fiddly optimisation problems
- It can be solved in $O(m^3)$ time (it involves inverting matrices) (phew it is not NP-complete!)
- We will see that there is an equivalent dual problem which allows us to use the kernel trick

Convexity of Lasso

- Recall that in Lasso we are asked to minimise

$$E = \sum_{j=1}^m (\mathbf{w}^\top \mathbf{x}_j - y_j)^2 + \nu \sum_{i=1}^p |w_i|$$

- We can rewrite this (using $w_i = w_i^+ - w_i^-$) as

$$E = \sum_{n=1}^N ((\mathbf{w}^+ - \mathbf{w}^-)^\top \mathbf{x}_n - y_n)^2 + \nu \sum_{i=1}^p (w_i^+ + w_i^-)$$

subject to $w_i^+, w_i^- \geq 0$ for all i

- Again it Lasso has a unique minima